



Deep Learning

Credit Card Fraud Detection

U Kang
Seoul National University



In This Lecture

- Credit card transactions
- Fraud detection
- Imbalanced classification
- Feedforward neural networks



Outline

- ➡ ☐ **Problem Definition**
- ☐ Requirements
- ☐ Preprocessing Codes
- ☐ Answers
- ☐ Autoencoder



Dataset (1)

- Transactions by credit cards in two days
- Each transaction contains a binary class
 - Normal (class 0)
 - Fraud (class 1)
- Most features are numerical values
 - They are the result of a PCA transformation
 - No information due to confidentiality issues



Dataset (2)

- Contains 28 numerical features (V1 ~ V28)
- Contains 2 special features ('Time' and 'Amount')
 - 'Time' is the seconds from the start
 - 'Amount' is the transaction amount
 - They are not the result of PCA
- The column 'Class' is the response variable
 - 1 in case of fraud 0 otherwise



Problem Definition

- To classify transactions into **Fraud** or **Normal**
- Simple classification problem on binary classes
- NOTE: the dataset is highly imbalanced!
 - 492 frauds out of 284,807 transactions (0.172%)



Selection of an Algorithm

- The dataset has no meaningful structures
 - Images, time-series, ...
- Thus, we use feedforward neural networks



Outline

☒ Problem Definition

 ☐ **Requirements**

☐ Preprocessing Codes

☐ Answers

☐ Autoencoder



Considering Imbalance (1)

- We need to consider the imbalance
 - Otherwise, the model will produce only 'Normal'
- Two famous techniques when training:
 - Undersampling
 - Select a subset of normal transactions
 - Oversampling
 - Duplicate fraud transactions
 - Adjust weights of the transactions



Considering Imbalance (2)

- Is *accuracy* enough for the evaluation?
- If we classify all instances as 'Normal', the accuracy will become high
- But, that is definitely not we want
- We need another evaluation metric



Model Structure

- Implement 3 dense layers
 - Numbers of hidden nodes: (64, 32, 16)
- Use dropout only in the third hidden layer
- Activation function: ReLU
- Optimizer: Adam
- Loss function: cross entropy



Feature Distribution

- We have the two categories of features
 - Simple numerical features (V_1, \dots, V_{28})
 - Special features (Time and Amount)
- They are distributed differently
 - Averages, standard deviations, ...
- We need to standardize them!



Outline

☒ Problem Definition

☒ Requirements

 ☐ **Preprocessing Codes**

☐ Answers

☐ Autoencoder



Loading the Dataset (1)

- Read the dataset using the pandas package:

```
import pandas as pd
import tensorflow as tf

df = pd.read_csv('creditcard.csv')
```

- The variable df is a DataFrame instance
- It is a 2D table of instances and features



Loading the Dataset (2)

- You can check the contents:

```
df.head()
```

	Time	V1	V2	V3	V4	V5	V6
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921



Dividing the Dataset (1)

- Divide the dataset by the class:

```
fraud = df[df.Class == 1]  
normal = df[df.Class == 0]
```

- Here [df.Class == 1] selects the normal ones
- This is row indexing (cf. column indexing)



Dividing the Dataset (2)

- Create training and test sets:

```
X_train = pd.concat([fraud.sample(frac=0.8,  
                                random_state=0),  
                    normal.sample(frac=0.8,  
                                random_state=0)],  
                    axis=0)  
X_test = df.loc[~df.index.isin(X_train.index)]
```

- We divide it with the ratio of 8:2
- We consider each class separately (why?)



Dividing the Dataset (3)

- Shuffle the resulting instances:

```
from sklearn.utils import shuffle  
  
X_train = shuffle(X_train, random_state=0)  
X_test = shuffle(X_test, random_state=0)
```

- This is useful for most ML algorithms



Creating Labels

- Create the labels for training and test sets:

```
y_train = X_train['Class'].values.astype('float')
y_train = np.stack([1 - y_train, y_train]).T

y_test = X_test['Class'].values.astype('float')
y_test = np.stack([1 - y_test, y_test]).T
```

- Note that each label is one-hot encoded



Deleting the Labels

- Delete the labels from the DataFrames:

```
fields = ['Class']  
X_train = X_train.drop(fields, axis=1)  
X_test = X_test.drop(fields, axis=1)
```

- We have 30 features currently
- Is it fine to use 'Amount' and 'Time' as it is?



Questions?