



Data Intelligence

Recommendation-1

Content based & Collaborative Filtering

U Kang
Seoul National University



In This Lecture

- Understand the motivation and the problem of recommendation
- Compare the content-based vs. collaborative filtering approaches for recommender system
- Learn how to evaluate methods for recommendation




Outline

- ➡ ☐ **Overview**
- ☐ Content-based Recommender System
- ☐ Collaborative Filtering
- ☐ Evaluation & Complexity



Example: Recommender Systems



■ Customer X

- ❑ Buys Metallica CD
- ❑ Buys Megadeth CD

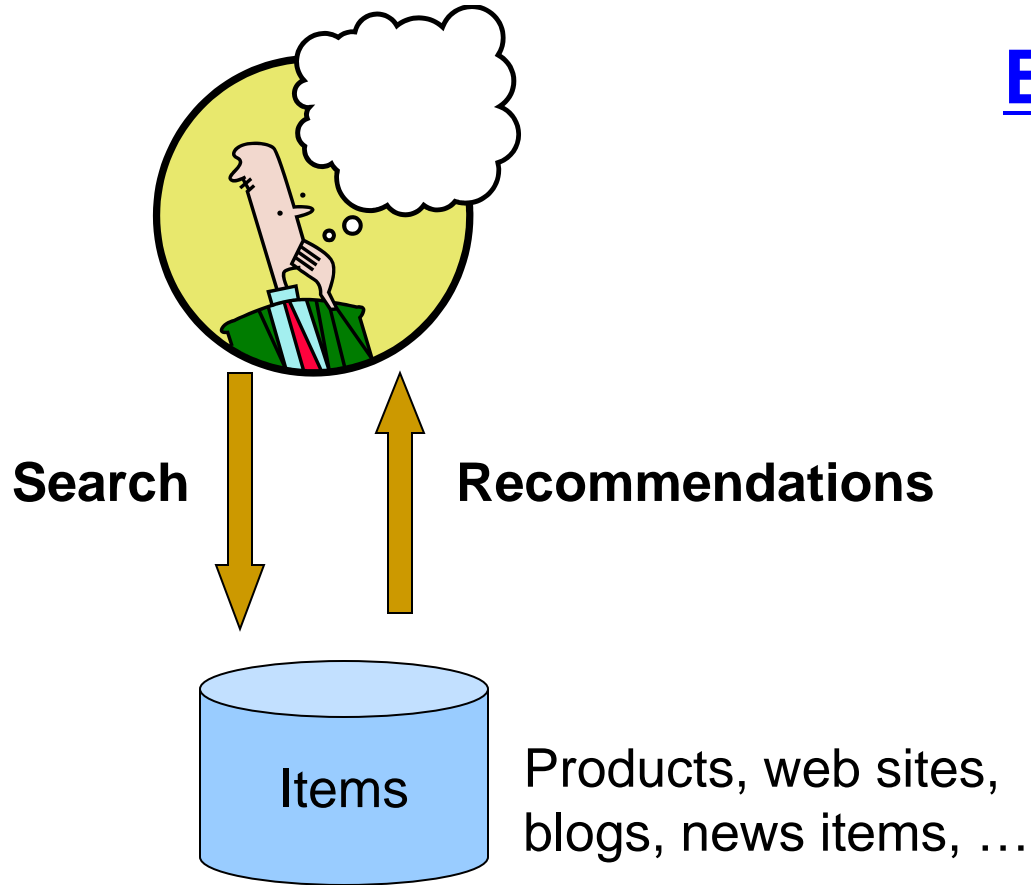


■ Customer Y

- ❑ Does search on Metallica
- ❑ Recommender system suggests Megadeth from data collected about customer X



Recommendations



Examples:

amazon.com



StumbleUpon



del.icio.us



movielens

helping you find the *right* movies

last.fm
the social music revolution

Google
News

YouTube

XBOX
LIVE



Offline vs. Online Recommendation

- **Offline recommendation: popular item**
 - ❑ Wall-mart: shelf space contains only 'popular' items
 - ❑ Also: TV networks, movie theaters,...
- **Web enables near-zero-cost dissemination of information about products**
 - ❑ Can recommend scarce items, too
- **More choice necessitates better filters**
 - ❑ Recommendation engines
 - ❑ How **Into Thin Air (1998)** made **Touching the Void (1988)** a bestseller: <http://www.wired.com/wired/archive/12.10/tail.html>



Sidenote: The Long Tail



Sources: Erik Brynjolfsson and Jeffrey Hu, MIT, and Michael Smith, Carnegie Mellon; Barnes & Noble; Netflix; RealNetworks
Source: Chris Anderson (2004)



Types of Recommendations

- **Editorial and hand curated**

- List of favorite cities
- List of “essential” items for travel

- **Simple aggregates**

- Top 10, Most Popular, Recent Uploads

- **Tailored to individual users**

- Amazon, Netflix, ...



Formal Model

- X = set of **Customers**
- S = set of **Items**
- **Utility function** $u: X \times S \rightarrow R$
 - R = set of ratings
 - R is a totally ordered set
 - e.g., **0-5** stars, real number in **[0,1]**



Utility Matrix

| | Avatar | LOTR | Matrix | Pirates |
|-------|--------|------|--------|---------|
| Alice | 1 | | 0.2 | |
| Bob | | 0.5 | | 0.3 |
| Carol | 0.2 | | 1 | |
| David | | | | 0.4 |



Key Problems

- **(1) Gathering “known” ratings for matrix**
 - How to collect the data in the utility matrix
- **(2) Extrapolate unknown ratings from the known ones**
 - Mainly interested in high unknown ratings
 - We are not interested in knowing what you don't like but what you like
- **(3) Evaluating extrapolation methods**
 - How to measure success/performance of recommendation methods



(1) Gathering Ratings

■ Explicit

- ❑ Ask people to rate items
- ❑ Doesn't work well in practice – people can't be bothered

■ Implicit

- ❑ Learn ratings from user actions
 - E.g., purchase implies high rating
- ❑ What about low ratings?
 - “not buying an item” = “don't like the item” ?



(2) Extrapolating Utilities

- **Key problem:** Utility matrix U is **sparse**
 - Most people have not rated most items
 - **Cold start:**
 - New items have no ratings
 - New users have no history
- **Three approaches to recommender systems:**
 - **1)** Content-based
 - **2)** Collaborative
 - **3)** Latent factor based



Outline

☒ Overview

 ☐ **Content-based Recommender System**

☐ Collaborative Filtering

☐ Evaluation & Complexity



Content-based Recommendations

- **Main idea:** Recommend items to customer x similar to previous items rated highly by x
 - John enjoyed watching “Avengers Infinity War”. John will also like “Avengers End Game” as well since they are similar in content

Example:

- **Movie recommendations**

- Recommend movies with same actor(s), genre, ...

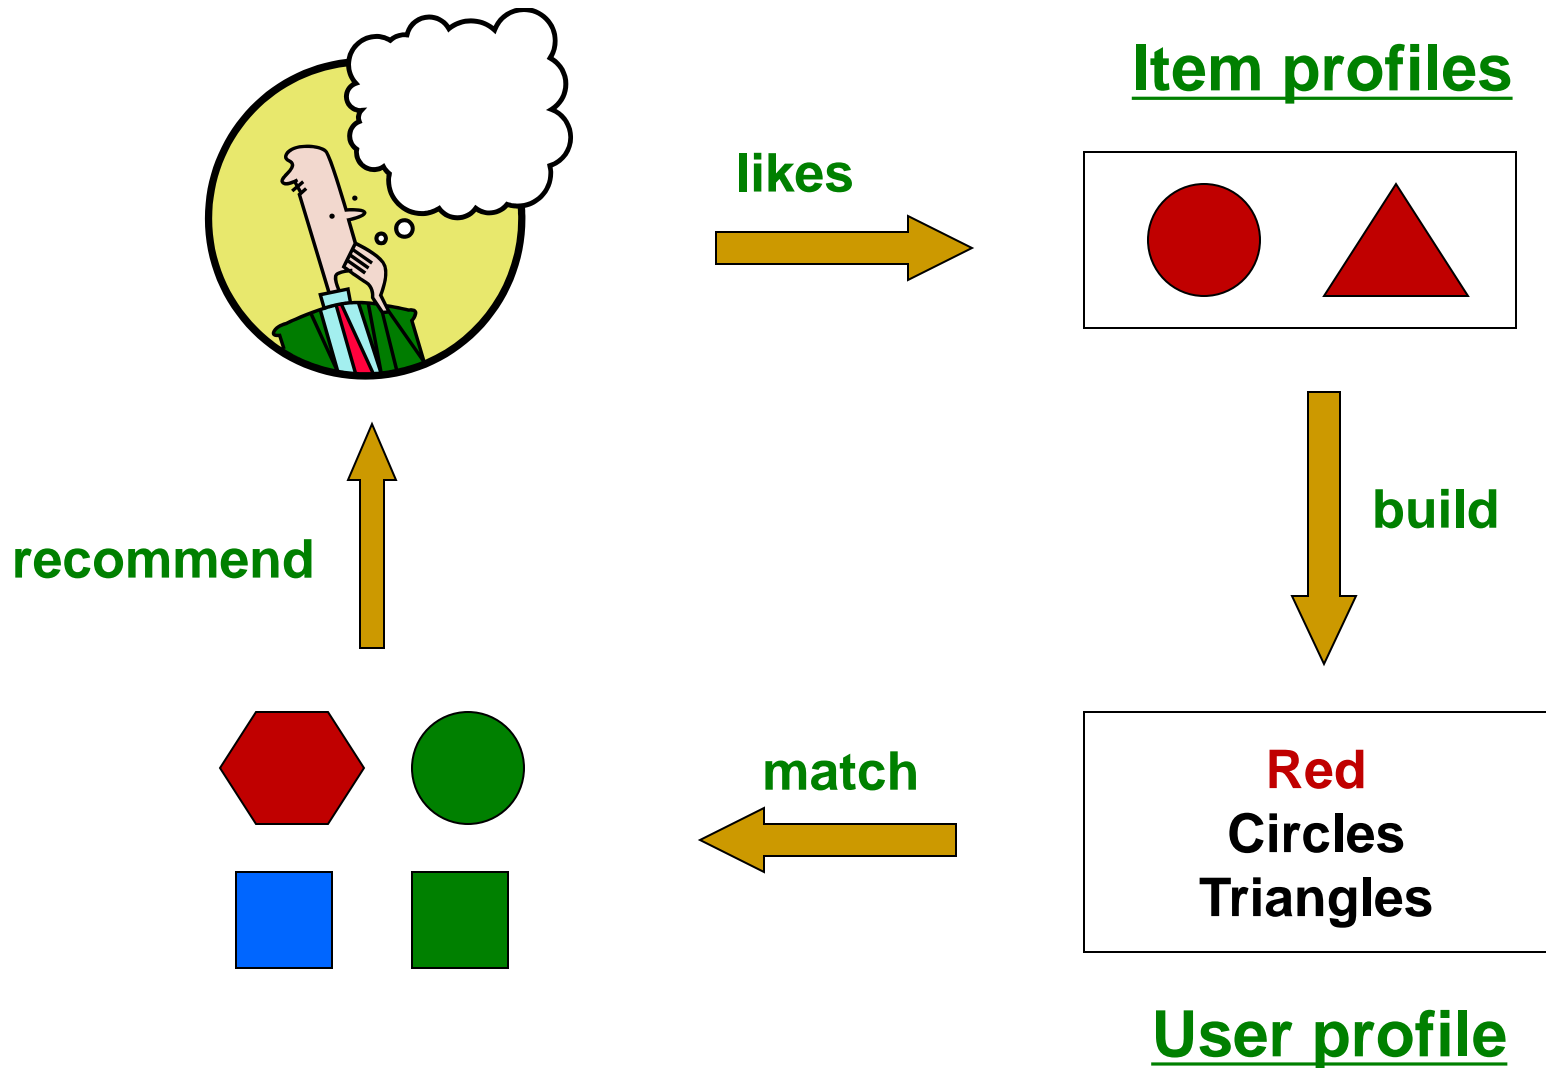
- **Websites, blogs, news**

- Recommend other sites with “similar” content





Plan of Action





Item Profiles

- For each item, create an **item profile**
- **Profile is a set (vector) of features**
 - **Movies:** author, title, actor, ...
 - **Text:** set of “important” words in document
- **How to pick important features?**
 - Usual heuristic from text mining is **TF-IDF**
(Term frequency * Inverse Doc Frequency)
 - **Term ... Feature**
 - **Document ... Item**



Sidenote: TF-IDF

f_{ij} = frequency of term (feature) i in doc (item) j

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}}$$

Note: we normalize TF to discount for “longer” documents

n_i = number of docs that mention term i

N = total number of docs

$$IDF_i = \log \frac{N}{n_i}$$

TF-IDF score: $w_{ij} = TF_{ij} \times IDF_i$

Doc profile = set of words with highest **TF-IDF** scores, together with their scores



User Profiles and Prediction

■ User profile possibilities:

- Weighted average of rated item profiles
- **Variation:** weight by difference from average rating for item
- ...

■ Prediction heuristic:

- Given user profile \mathbf{x} and item profile \mathbf{i} , estimate

$$u(\mathbf{x}, \mathbf{i}) = \cos(\mathbf{x}, \mathbf{i}) = \frac{\mathbf{x} \cdot \mathbf{i}}{||\mathbf{x}|| \cdot ||\mathbf{i}||}$$



Pros: Content-based Approach

- **+: No need for data on other users**
 - No cold-start or sparsity problems
- **+: Able to recommend to users with unique tastes**
- **+: Able to recommend new & unpopular items**
 - No first-rater problem
- **+: Able to provide explanations**
 - Can provide explanations of recommended items by listing content-features that caused an item to be recommended




Cons: Content-based Approach

- **–: Finding the appropriate features is hard**
 - E.g., images, movies, music
- **–: Recommendations for new users**
 - **How to build a user profile?**
- **–: Overspecialization**
 - Never recommends items outside user's content profile
 - Users want to be surprised sometimes
 - People might have multiple interests
 - **Unable to exploit quality judgments of other users**



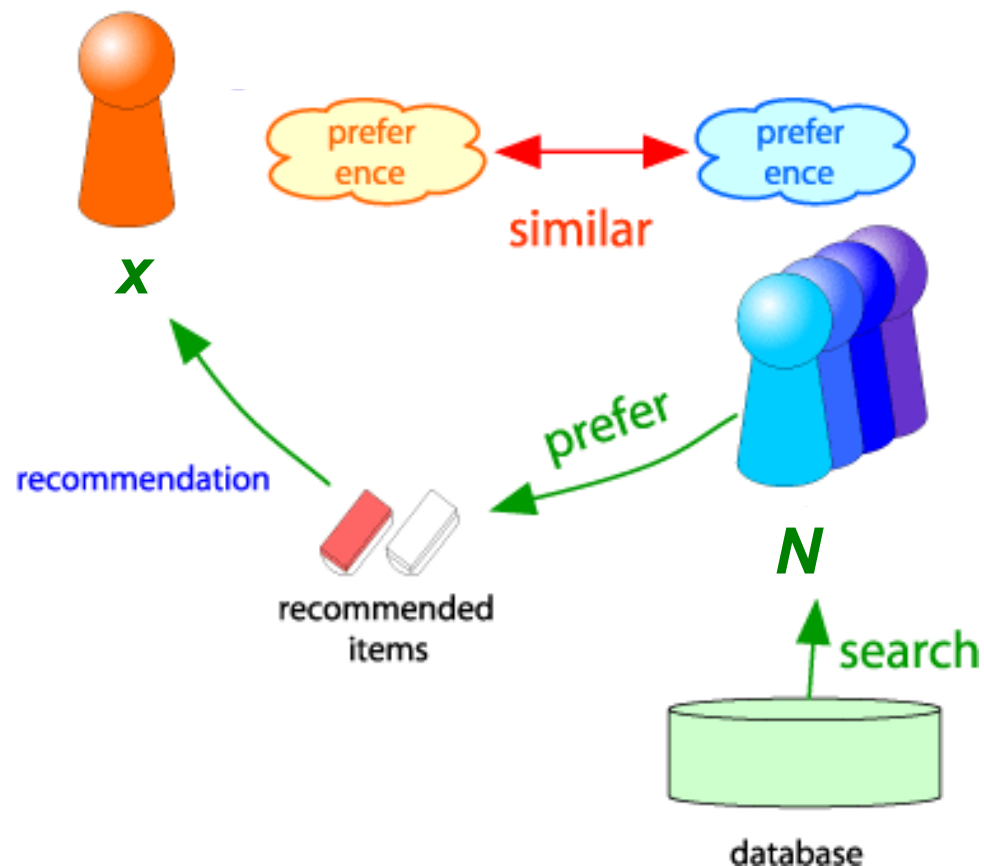
Outline

- ☒ Overview
- ☒ Content-based Recommender System
-  ☐ **Collaborative Filtering**
- ☐ Evaluation & Complexity



Collaborative Filtering

- Consider user x
- Find set N of other users whose ratings are “**similar**” to x ’s ratings
- Estimate x ’s ratings based on ratings of users in N



Note that contents of items are not used here.



Finding “Similar” Users

$$\begin{aligned} r_x &= [* , _ , _ , * , ***] \\ r_y &= [* , _ , ** , ** , _] \end{aligned}$$

- Let r_x be the vector of user x 's ratings

- **Jaccard similarity measure**

- **Problem:** Ignores the value of the rating

r_x, r_y as sets:

$$r_x = \{1, 4, 5\}$$

$$r_y = \{1, 3, 4\}$$

- **Cosine similarity measure**

- $\text{sim}(x, y) = \cos(r_x, r_y) = \frac{r_x \cdot r_y}{||r_x|| \cdot ||r_y||}$

- **Problem:** low rating is not penalized much

r_x, r_y as points:

$$r_x = \{1, 0, 0, 1, 3\}$$

$$r_y = \{1, 0, 2, 2, 0\}$$

- **Pearson correlation coefficient**

- S_{xy} = items rated by both users x and y

$$\text{sim}(x, y) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2} \sqrt{\sum_{s \in S_{xy}} (r_{ys} - \bar{r}_y)^2}}$$

$\bar{r}_x, \bar{r}_y \dots$ avg.
rating of x, y



Similarity Metric

| | HP1 | HP2 | HP3 | TW | SW1 | SW2 | SW3 |
|---|-----|-----|-----|----|-----|-----|-----|
| A | 4 | | | 5 | 1 | | |
| B | 5 | 5 | 4 | | | | |
| C | | | | 2 | 4 | 5 | |
| D | | 3 | | | | | 3 |

■ **Intuitively we want:** $\text{sim}(A, B) > \text{sim}(A, C)$

■ **Jaccard similarity:** $1/5 < 2/4$

■ **Cosine similarity:** $0.380 > 0.322$

□ Problem: low rating is not penalized much

□ **Solution: subtract the (row) mean**

| | HP1 | HP2 | HP3 | TW | SW1 | SW2 | SW3 |
|---|-----|-----|------|------|------|-----|-----|
| A | 2/3 | | | 5/3 | -7/3 | | |
| B | 1/3 | 1/3 | -2/3 | | | | |
| C | | | | -5/3 | 1/3 | 4/3 | |
| D | | 0 | | | | | 0 |

sim A,B vs. A,C:
 $0.092 > -0.559$



Rating Predictions

From similarity metric to recommendations:

- Let r_x be the vector of user x 's ratings
- Let N (called 'k-nearest neighbors') be the set of k users most similar to x who have rated item i
- **Prediction r_{xi} for item i of user x :**
 - $r_{xi} = \frac{1}{k} \sum_{y \in N} r_{yi}$
 - $r_{xi} = \frac{\sum_{y \in N} s_{xy} \cdot r_{yi}}{\sum_{y \in N} s_{xy}}$
- **Many other tricks possible...**

Shorthand:

$$s_{xy} = \text{sim}(x, y)$$



Item-Item Collaborative Filtering

- So far: **User-user collaborative filtering**
- **Another view: Item-item**
 - For item i , find other similar items rated by user x
 - Use the utility matrix for computing similarity
 - Estimate rating for item i based on ratings for similar items
 - Can use same similarity metrics and prediction functions as in user-user model

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

s_{ij} ... similarity of items i and j
 r_{xj} ... rating of user x on item j
 $N(i;x)$... set items rated by x similar to i




Item-Item CF ($|N|=2$)

| | | users | | | | | | | | | | | |
|--------|---|-------|---|---|---|---|---|---|---|---|----|----|----|
| | | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| movies | 1 | 1 | | 3 | | | 5 | | | 5 | | 4 | |
| | 2 | | | 5 | 4 | | | 4 | | | 2 | 1 | 3 |
| | 3 | 2 | 4 | | 1 | 2 | | 3 | | 4 | 3 | 5 | |
| | 4 | | 2 | 4 | | 5 | | | 4 | | | 2 | |
| | 5 | | | 4 | 3 | 4 | 2 | | | | | 2 | 5 |
| | 6 | 1 | | 3 | | 3 | | | 2 | | | 4 | |



- unknown rating



- rating between 1 to 5



Item-Item CF ($|N|=2$)

users

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|
| 1 | 1 | | 3 | | ? | 5 | | | 5 | | 4 | |
| 2 | | | 5 | 4 | | | 4 | | | 2 | 1 | 3 |
| 3 | 2 | 4 | | 1 | 2 | | 3 | | 4 | 3 | 5 | |
| 4 | | 2 | 4 | | 5 | | | 4 | | | 2 | |
| 5 | | | 4 | 3 | 4 | 2 | | | | | 2 | 5 |
| 6 | 1 | | 3 | | 3 | | | 2 | | | 4 | |

movies



- estimate rating of movie **1** by user **5**



Item-Item CF ($|N|=2$)

| | | users | | | | | | | | | | | | | |
|--------|----------|-------|---|---|---|---|---|---|---|---|----|----|----|-------------|-------|
| | | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | | |
| movies | 1 | 1 | | 3 | | ? | 5 | | | 5 | | 4 | | sim(1,m) | 1.00 |
| | 2 | | | 5 | 4 | | | 4 | | | 2 | 1 | 3 | | -0.18 |
| | <u>3</u> | 2 | 4 | | 1 | 2 | | 3 | | 4 | 3 | 5 | | <u>0.41</u> | |
| | 4 | | 2 | 4 | | 5 | | | 4 | | | 2 | | -0.10 | |
| | 5 | | | 4 | 3 | 4 | 2 | | | | | 2 | 5 | -0.31 | |
| | <u>6</u> | 1 | | 3 | | 3 | | | 2 | | | 4 | | <u>0.59</u> | |

Neighbor selection:

Identify movies similar to movie 1, rated by user 5

Similarity computation:

- 1) Subtract mean rating m_i from each movie i
 $m_1 = (1+3+5+5+4)/5 = 3.6$
row 1: [-2.6, 0, -0.6, 0, 0, 1.4, 0, 0, 1.4, 0, 0.4, 0]
- 2) Compute cosine similarities between rows



Item-Item CF ($|N|=2$)

| | | users | | | | | | | | | | | | sim(1,m) |
|--------|----------|-------|---|---|---|---|---|---|---|---|----|----|----|-------------|
| | | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | |
| movies | 1 | 1 | | 3 | | ? | 5 | | | 5 | | 4 | | 1.00 |
| | 2 | | | 5 | 4 | | | 4 | | | 2 | 1 | 3 | -0.18 |
| | <u>3</u> | 2 | 4 | | 1 | 2 | | 3 | | 4 | 3 | 5 | | <u>0.41</u> |
| | 4 | | 2 | 4 | | 5 | | | 4 | | | 2 | | -0.10 |
| | 5 | | | 4 | 3 | 4 | 2 | | | | | 2 | 5 | -0.31 |
| | <u>6</u> | 1 | | 3 | | 3 | | | 2 | | | 4 | | <u>0.59</u> |

Compute similarity weights:

$s_{1,3}=0.41$, $s_{1,6}=0.59$



Item-Item CF ($|N|=2$)

| | | users | | | | | | | | | | | |
|--------|----------|-------|---|---|---|-----|---|---|---|---|----|----|----|
| | | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
| movies | 1 | 1 | | 3 | | 2.6 | 5 | | | 5 | | 4 | |
| | 2 | | | 5 | 4 | | | 4 | | | 2 | 1 | 3 |
| | <u>3</u> | 2 | 4 | | 1 | 2 | | 3 | | 4 | 3 | 5 | |
| | 4 | | 2 | 4 | | 5 | | | 4 | | | 2 | |
| | 5 | | | 4 | 3 | 4 | 2 | | | | | 2 | 5 |
| | <u>6</u> | 1 | | 3 | | 3 | | | 2 | | | 4 | |

Predict by taking weighted average:

$$r_{1.5} = (0.41 \cdot 2 + 0.59 \cdot 3) / (0.41 + 0.59) = 2.6$$

$$r_{ix} = \frac{\sum_{j \in N(i;x)} s_{ij} \cdot r_{jx}}{\sum s_{ij}}$$



CF: Common Practice

Before:

$$r_{xi} = \frac{\sum_{j \in N(i;x)} s_{ij} r_{xj}}{\sum_{j \in N(i;x)} s_{ij}}$$

- Define **similarity** s_{ij} of items i and j
- Select k nearest neighbors $N(i; x)$
 - Items most similar to i , that were rated by x
- Estimate rating r_{xi} as the weighted average:

$$r_{xi} = b_{xi} + \frac{\sum_{j \in N(i;x)} s_{ij} \cdot (r_{xj} - b_{xj})}{\sum_{j \in N(i;x)} s_{ij}}$$

baseline estimate for r_{xi}

$$b_{xi} = \mu + b_x + b_i$$

- μ = overall mean movie rating
- b_x = rating deviation of user x
= (avg. rating of user x) - μ
- b_i = rating deviation of movie i
= (avg. rating of movie i) - μ



CF: Baseline Predictor

- Mean movie rating: **3.7 stars**
- *The Sixth Sense* is **0.5** stars above avg.
- Joe rates **0.2** stars below avg.

⇒ **Baseline estimation:**

Joe will rate *The Sixth Sense* 4 stars





Item-Item vs. User-User

| | Avatar | LOTR | Matrix | Pirates |
|-------|--------|------|--------|---------|
| Alice | 1 | | 0.8 | |
| Bob | | 0.5 | | 0.3 |
| Carol | 0.9 | | 1 | 0.8 |
| David | | | 1 | 0.4 |

- In practice, it has been observed that item-item often works better than user-user
- **Why?** Items are simpler, users have multiple tastes



Pros/Cons of Collaborative Filtering

- **+ Works for any kind of item**
 - No feature selection needed
- **+ Can use other people's suggestions**



Pros/Cons of Collaborative Filtering

■ - Cold Start:

- ❑ Needs enough users in the system to find a match

■ - Sparsity:

- ❑ The user/ratings matrix is sparse
- ❑ Hard to find users that have rated the same items

■ - First rater:

- ❑ Cannot recommend an item that has not been previously rated (e.g., new items, esoteric items)

■ - Popularity bias:

- ❑ Cannot recommend items to someone with unique taste
- ❑ Tends to recommend popular items




Hybrid Methods

- **Implement two or more different recommenders and combine predictions**
 - Perhaps using a linear model
- **Add content-based methods to collaborative filtering**
 - Item profiles for new item problem
 - Demographics to deal with new user problem



Outline

- ☒ Overview
- ☒ Content-based Recommender System
- ☒ Collaborative Filtering
-  ☐ **Evaluation & Complexity**



Evaluation

movies

users

| | | | | | |
|---|---|---|---|---|---|
| 1 | 3 | 4 | | | |
| | 3 | 5 | | | 5 |
| | | 4 | 5 | | 5 |
| | | 3 | | | |
| | | 3 | | | |
| 2 | | | 2 | | 2 |
| | | | | 5 | |
| | 2 | 1 | | | 1 |
| | 3 | | | 3 | |
| 1 | | | | | |



Evaluation

movies

users

| | | | | | |
|---|---|---|---|---|---|
| 1 | 3 | 4 | | | |
| | 3 | 5 | | | 5 |
| | | 4 | 5 | | 5 |
| | | 3 | | | |
| | | 3 | | | |
| 2 | | | ? | | ? |
| | | | | ? | |
| | 2 | 1 | | | ? |
| | 3 | | | ? | |
| 1 | | | | | |

Test Data Set



Evaluating Predictions

■ Compare predictions with known ratings

□ Root-mean-square error (RMSE)


- $\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$ where r_{xi} is predicted, r_{xi}^* is the true rating of x on i

□ Precision at top 10: error in top 10 highest predictions

□ Rank Correlation:


- Spearman's *correlation* between system's and user's complete rankings

(From Wikipedia)



Pearson correlation coefficient

U Kang



Rank correlation coefficient=1



Problems with Error Measures

- **Narrow focus on accuracy sometimes misses the point**
 - E.g., prediction diversity
- **In practice, we care only to predict high ratings:**
 - RMSE might penalize a method that does well for high ratings and badly for others



Collaborative Filtering: Complexity

- Expensive step is finding k most similar customers: $O(|X|)$
 - X ... set of customers
- **Too expensive to do at runtime**
 - Could pre-compute
- Pre-compute finding similar customers
 - Near-neighbor search in high dimensions (**LSH**)
 - Clustering
 - Dimensionality reduction



Tip: Add Data

- **Simple method on large data is better than complex method on small data**

- Leverage all the data
- Don't try to reduce data size in an effort to make fancy algorithms work

- **Add more data**

- e.g., add IMDB data on genres

- **More data beats better algorithms**

<http://anand.typepad.com/datawocky/2008/03/more-data-usual.html>



What You Need to Know

- Motivation and the problem of recommendation
- Compare the content-based vs. collaborative filtering approaches for recommender system
 - Content-based: less cold-start problem
 - Collaborative filtering: works for any item
- Evaluation methods for recommendation
 - Training set and test set



Questions?