



Data Intelligence

Deep Matrix Factorization - Lab

U Kang
Seoul National University



In This Lecture

- Deep Matrix Factorization
 - Recommendation
 - Deep Learning based matrix factorization
 - Data preparation
 - Model implementation
 - Model training / evaluation

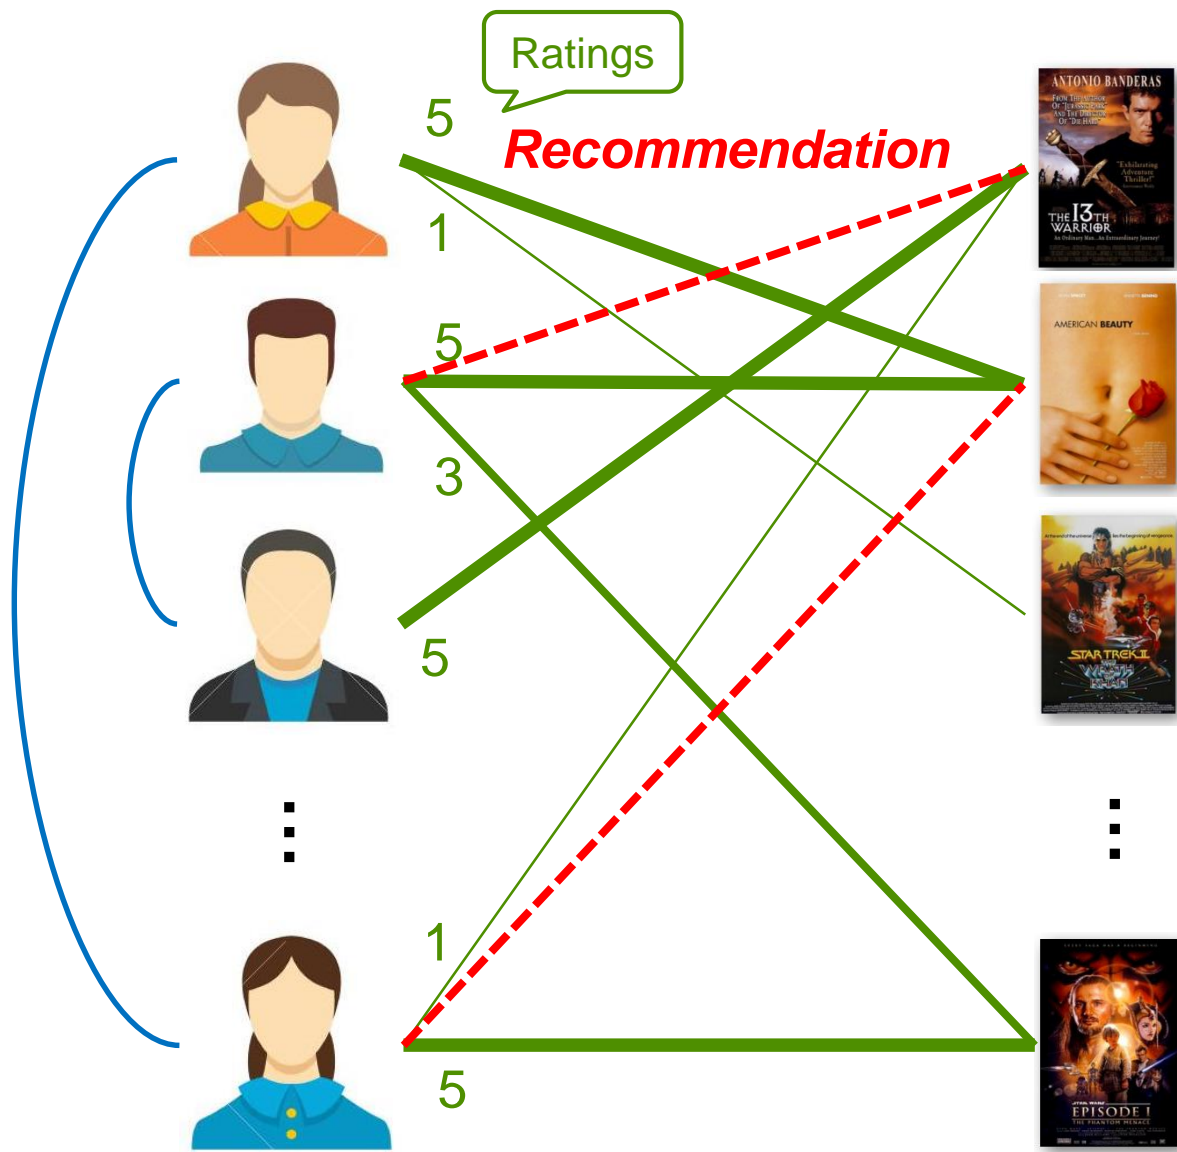


Outline

- ➡ ☐ **Recommendation**
 - ☐ Deep Learning based Matrix Factorization
 - ☐ Data Preparation
 - ☐ Model Implementation
 - ☐ Model Training / Evaluation



Recommendation





Matrix Completion

- Matrix Completion is a surrogate problem of recommendation
 - Users want to be provided items that they will give high ratings
- Matrix Completion
 - **Given:** a sparse rating matrix R
 - **Goal:** to predict unseen rating values in R

R matrix

User \ Item	Item				
	1	2	3	4	5
1	5	?	?	2	1
2	?	?	2	?	?
3	3	?	?	?	?
4	?	3	1	?	?
5	5	?	?	2	?

User-Movie Matrix



Outline

☒ Recommendation

 ☐ **Deep Learning based Matrix Factorization**

☐ Data Preparation

☐ Model Implementation

☐ Model Training / Evaluation



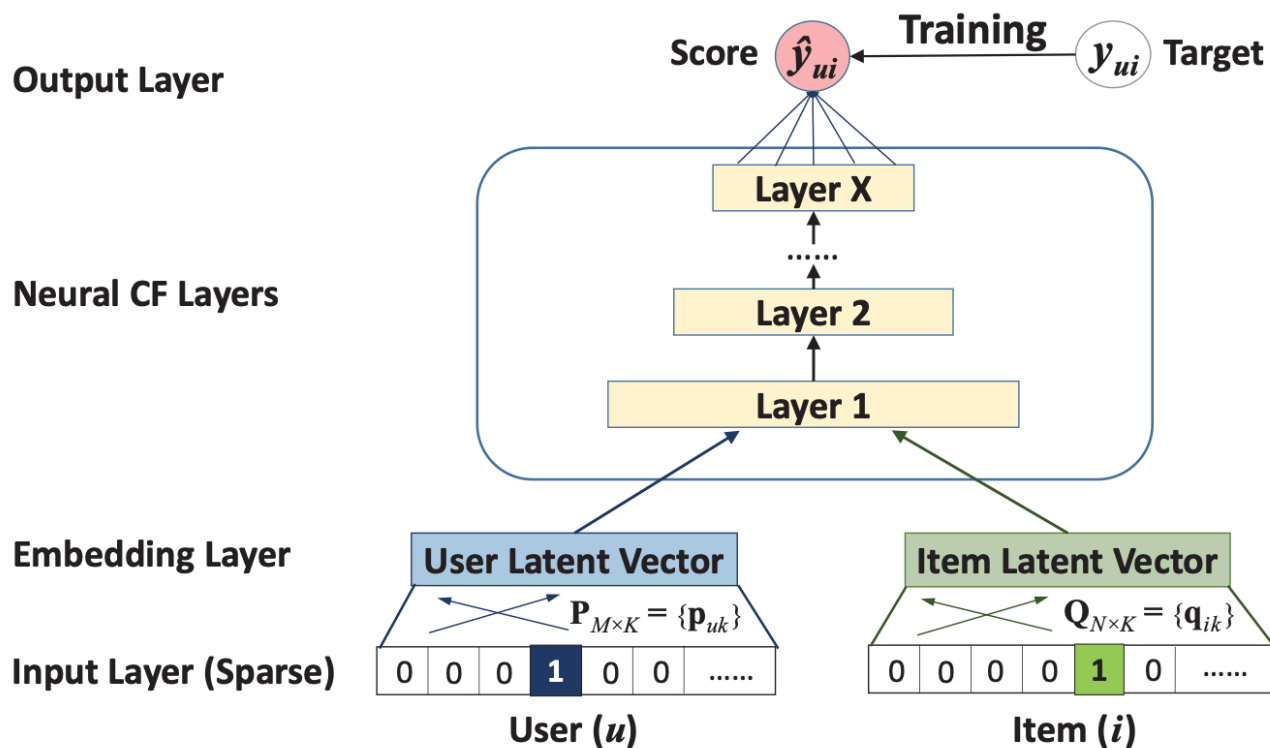
Matrix Factorization

- Limitations of matrix factorization
 - A linear model
 - It cannot model non-linear relationships
 - Limited parameters
 - Only user and item latent factors are trained
 - Simple modeling
 - It expresses user-item relationships by the inner-product which is too simple to model complicated real-world interactions



Deep Matrix Factorization (1)

■ Deep learning based matrix factorization



He, Xiangnan, et al. "Neural collaborative filtering." WWW. 2017.




Deep Matrix Factorization (2)

- Deep learning based matrix factorization
 - A non-linear model
 - It models non-linear relationships
 - Wide range of parameters
 - It learns not only user and item latent factors, but also user-item relationships
 - Complex modeling
 - It can express complicated user-item relationships through deep neural networks



Outline

- ☒ Recommendation
- ☒ Deep Learning based Matrix Factorization
-  ☐ **Data Preparation**
- ☐ Model Implementation
- ☐ Model Training / Evaluation



PyTorch

- PyTorch (<https://pytorch.org/>)
 - A deep learning framework
 - It supports GPU computations
 - Model training (back-propagation) can be done easily
 - Numpy friendly
 - The usage is similar to Numpy
 - Includes a lot of useful functions
 - Dataloader
 - Model save/load
 - Initialization
 - Visualization
 - ...



Reading Data File (1)

- Set the data path and split ratio
 - “ratings.dat” contains interaction logs

```
import numpy as np
```

```
in_path = './data/ml-1m-raw/'  
rating_file = in_path + 'ratings.dat'
```



Reading Data File (2)

- Read data file
 - Format of “ratings.dat”
 - user_id::item_id::rating::time_stamp

```
raw = []  
with open(rating_file, 'r') as f_read:  
    for line in f_read.readlines():  
        line_list = line.split('::')  
        raw.append(line_list)  
raw = np.array(raw, dtype=np.int)
```



Assign New IDs

- We need new ids that start from 0

```
user_ids = list()
item_ids = list()
user_map = dict() # raw -> new
item_map = dict() # raw -> new

user_ids = np.unique(raw[:, 0])
item_ids = np.unique(raw[:, 1])

user_map = {v: i for (i, v) in enumerate(user_ids)}
item_map = {v: i for (i, v) in enumerate(item_ids)}

new = [[user_map[u], item_map[i], r] for (u, i, r)
        in zip(raw[:, 0], raw[:, 1], raw[:, 2])] # new array
new = np.array(new)
print(new.shape)
```

(1000209, 3)



Split Dataset

- Randomly split the dataset into training/test sets

```
from sklearn.model_selection import train_test_split
```

```
train, test = train_test_split(new, test_size=0.2, shuffle=True, random_state=42)
```

```
print(train.shape)  
print(test.shape)
```

```
(800167, 3)
```

```
(200042, 3)
```



Dataset and Data Loader (1)

```
from torch.utils.data import Dataset, DataLoader
```

- PyTorch supports custom dataset and data loader
 - Dataset
 - We should define “__len__()” and “__getitem__()” functions
 - “__len__()”: Length of the dataset
 - “__getitem__()”: How to get data when an index is given
 - DataLoader
 - We should define batch size and whether to shuffle the data



Dataset and Data Loader (2)

■ Define custom dataset

```
class MovieLensDataset(Dataset):  
    def __init__(self, x, y):  
        self.x = x  
        self.y = y  
  
    def __len__(self):  
        return len(self.x)  
  
    def __getitem__(self, idx):  
        x = torch.LongTensor(self.x[idx, :])  
        y = torch.FloatTensor(self.y[idx, :])  
        return x, y
```

- ❑ **x** and **y** are given to the object
- ❑ `__len__()` returns the length of the dataset
- ❑ `__getitem__()` returns the **x** and **y** at the given **idx**



Dataset and Data Loader (3)

■ Make train/test datasets and data loaders


```
train_dataset = MovieLensDataset(train[:, :-1], np.expand_dims(train[:, -1], axis=1))  
train_dataloader = DataLoader(train_dataset, batch_size=512, shuffle=True)
```

```
test_dataset = MovieLensDataset(test[:, :-1], np.expand_dims(test[:, -1], axis=1))  
test_dataloader = DataLoader(test_dataset, batch_size=len(test_dataset), shuffle=False)
```

- The data loaders will yield a batch of size “batch_size” whenever they are called
- We set batch_size to 512



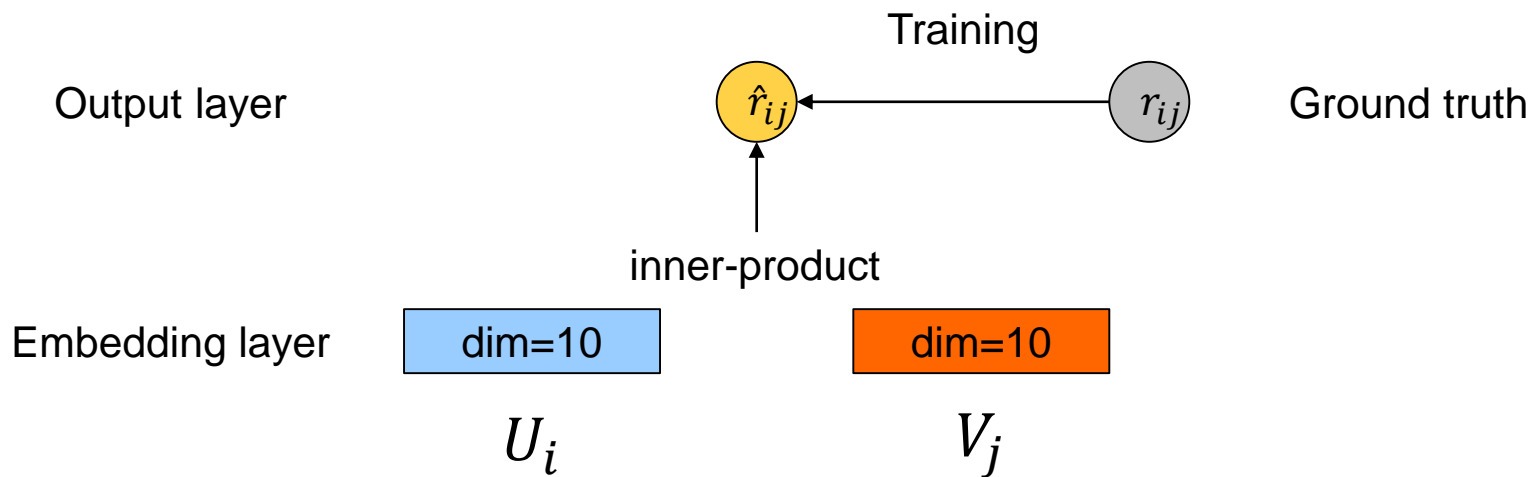
Outline

- ☒ Recommendation
- ☒ Deep Learning based Matrix Factorization
- ☒ Data Preparation
-  ☐ **Model Implementation**
- ☐ Model Training / Evaluation



Model (1)

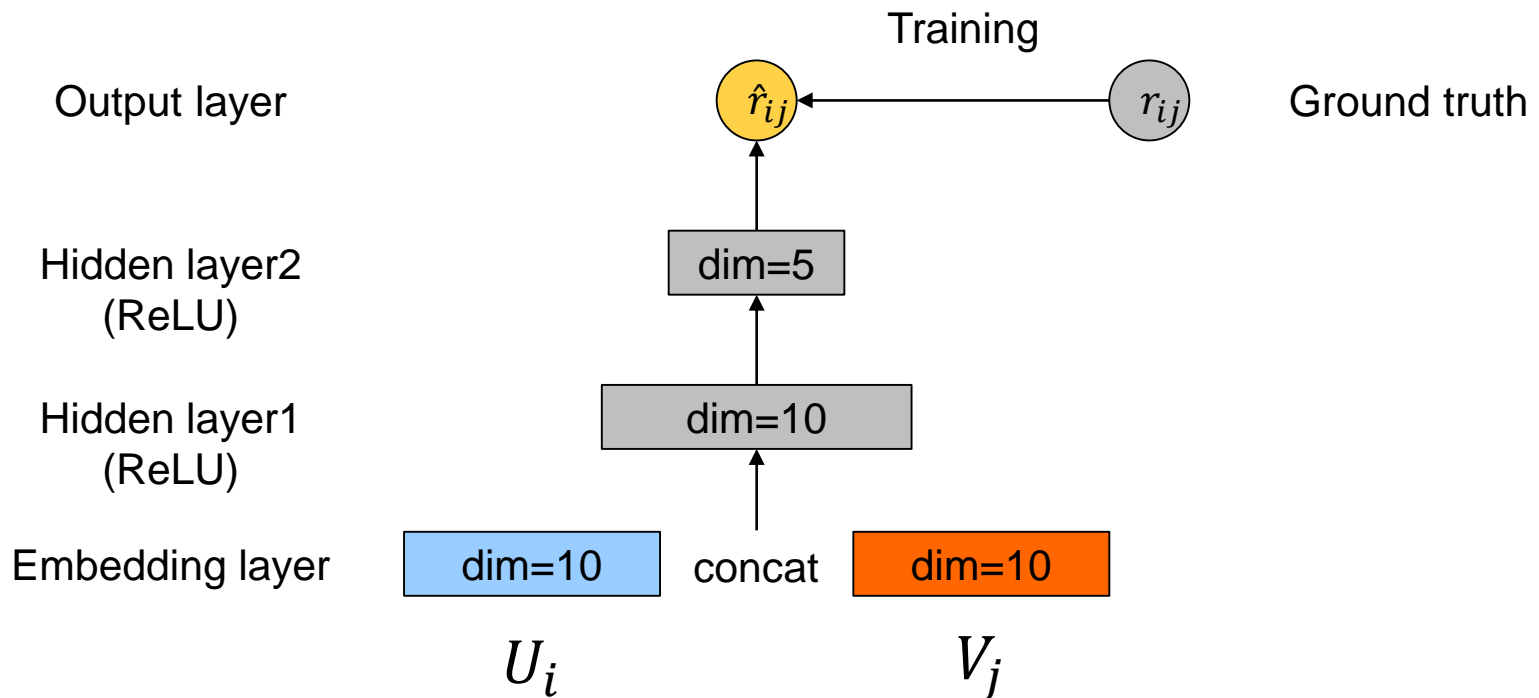
- We compare two models: MF and DeepMF
 - MF: a simple linear model





Model (2)

- We compare two models: MF and DeepMF
 - DeepMF: 3-layered fully-connected neural network





Hyper-parameters

■ Let's define hyper-parameters:

```
num_users = max(max(train[:, 0]), max(test[:, 0])) + 1  
num_items = max(max(test[:, 1]), max(test[:, 1])) + 1
```

```
K = 10  
hidden_dim1 = 10  
hidden_dim2 = 5  
lr = 1e-3  
decay = 1e-5  
epochs = 100
```

- ❑ K: embedding vector dimensionality
- ❑ hidden_dim: hidden layer dimensionality
- ❑ lr: learning rate
- ❑ decay: weight decay
- ❑ epochs: number of epochs



Import Library

■ Let's import libraries:

```
import torch
import time
DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

- ❑ We can choose whether to use GPU (CUDA)
- ❑ `torch.cuda.is_available()`: whether cuda is available



Define Models (1)

■ Let's define Matrix Factorization:

```
class MF(torch.nn.Module):
    def __init__(self, num_users, num_items, K):
        super().__init__()
        self.user_emb = torch.nn.Embedding(num_users, K)
        self.item_emb = torch.nn.Embedding(num_items, K)

    def forward(self, user_idx, item_idx):
        out = (self.user_emb(user_idx) * self.item_emb(item_idx)).sum(1, keepdim=True)
        return out
```

- ❑ torch.nn.Embedding: embedding parameters
- ❑ We should define “forward” function which is used in prediction



Define Models (2)

■ Let's define Deep Matrix Factorization:

```
class DeepMF(torch.nn.Module):
    def __init__(self, num_users, num_items, K, hidden_dim1, hidden_dim2):
        super().__init__()
        self.user_emb = torch.nn.Embedding(num_users, K)
        self.item_emb = torch.nn.Embedding(num_items, K)
        self.layer1 = torch.nn.Linear(2*K, hidden_dim1)
        self.layer2 = torch.nn.Linear(hidden_dim1, hidden_dim2)
        self.out = torch.nn.Linear(hidden_dim2, 1)
        self.activation = torch.nn.ReLU()

    def forward(self, user_idx, item_idx):
        out = torch.cat((self.user_emb(user_idx), self.item_emb(item_idx)), dim=1)
        out = self.activation(self.layer1(out))
        out = self.activation(self.layer2(out))
        out = self.out(out)
        return out
```

- ❑ torch.nn.Embedding: embedding parameters
- ❑ torch.nn.Linear: linear layer with bias
- ❑ torch.nn.ReLU: activation function



Make Models

- Make one of the two models: MF and DeepMF

```
model = MF(num_users, num_items, K)
# model = DeepMF(num_users, num_items, K, hidden_dim1, hidden_dim2)
model.to(DEVICE)
```

- torch.nn.Module class supports “to” function
 - It loads the parameters to CPU or GPU



Loss and Optimizer

```
criterion = torch.nn.MSELoss()  
optimizer = torch.optim.Adam(model.parameters(), lr=lr, weight_decay=decay)
```

- We use Mean Squared Error (MSE) as the loss

- $$MSE = \frac{\sum_u \sum_i (R_{ui} - \hat{R}_{ui})^2}{|ratings|}$$

- where

- R_{ui} is a ground-truth rating value
 - \hat{R}_{ui} is a predicted rating value by the model


- PyTorch supports the MSE loss

- We use Adam optimizer

- Parameters should be given to the optimizer



Outline

- ☒ Recommendation
- ☒ Deep Learning based Matrix Factorization
- ☒ Data Preparation
- ☒ Model Implementation
-  ☐ **Model Training / Evaluation**



Training and Evaluation

- In every epoch, train the model and evaluate it:

```
for epoch in range(epochs):  
    start_time = time.time()  
    train_mse = 0.  
    test_mse = 0.  
  
    # train the model  
  
    # test the model
```

- Training/testing codes are in the next slides



Training

- Train the model using the training dataset:

```
# train the model
model.train()
for batch_idx, (x, r) in enumerate(train_dataloader):
    # get data
    x, r = x.to(DEVICE), r.to(DEVICE)
    i, j = x[:, 0], x[:, 1]

    # set gradients to zero
    optimizer.zero_grad()

    # predict ratings
    pred = model(i, j)

    # get loss
    loss = criterion(pred, r)
    train_mse += loss.item()

    # backpropagation
    loss.backward()

    # update the parameters
    optimizer.step()

train_rmse = (train_mse/(batch_idx+1))**.5
```



Evaluation

- Test the model using the testing dataset:

```
# test the model
model.eval()
for batch_idx, (x, r) in enumerate(test_dataloader):
    # get data
    x, r = x.to(DEVICE), r.to(DEVICE)
    i, j = x[:, 0], x[:, 1]

    # predict ratings
    pred = model(i, j)

    # get loss
    loss = criterion(pred, r)
    test_mse += loss.item()

test_rmse = (test_mse/(batch_idx+1))**.5

end_time = time.time()
print(f'[{end_time-start_time:.2f}] Epoch: {epoch+1:3d}, '
      f'TrnRMSE: {train_rmse:.4f}, TestRMSE: {test_rmse:.4f}')
```



Results (1)

■ Matrix Factorization

[20.66]	Epoch:	1,	TrnRMSE:	4.6855,	TestRMSE:	4.4827
[20.70]	Epoch:	2,	TrnRMSE:	4.3197,	TestRMSE:	4.2276
[20.90]	Epoch:	3,	TrnRMSE:	4.0864,	TestRMSE:	4.0496
[20.68]	Epoch:	4,	TrnRMSE:	3.8859,	TestRMSE:	3.8329
[20.72]	Epoch:	5,	TrnRMSE:	3.5282,	TestRMSE:	3.2955
[20.57]	Epoch:	6,	TrnRMSE:	2.7649,	TestRMSE:	2.3986
[20.71]	Epoch:	7,	TrnRMSE:	1.9648,	TestRMSE:	1.7797
[20.88]	Epoch:	8,	TrnRMSE:	1.4992,	TestRMSE:	1.4422
[20.77]	Epoch:	9,	TrnRMSE:	1.2463,	TestRMSE:	1.2508
[20.84]	Epoch:	10,	TrnRMSE:	1.1043,	TestRMSE:	1.1392
[20.87]	Epoch:	11,	TrnRMSE:	1.0232,	TestRMSE:	1.0735
[20.62]	Epoch:	12,	TrnRMSE:	0.9758,	TestRMSE:	1.0334
[20.83]	Epoch:	13,	TrnRMSE:	0.9470,	TestRMSE:	1.0086
[21.35]	Epoch:	14,	TrnRMSE:	0.9284,	TestRMSE:	0.9921
[20.93]	Epoch:	15,	TrnRMSE:	0.9155,	TestRMSE:	0.9811
[20.82]	Epoch:	16,	TrnRMSE:	0.9056,	TestRMSE:	0.9725
[20.78]	Epoch:	17,	TrnRMSE:	0.8972,	TestRMSE:	0.9655
[20.65]	Epoch:	18,	TrnRMSE:	0.8898,	TestRMSE:	0.9595
[20.67]	Epoch:	19,	TrnRMSE:	0.8829,	TestRMSE:	0.9542
[20.83]	Epoch:	20,	TrnRMSE:	0.8765,	TestRMSE:	0.9497
[20.72]	Epoch:	21,	TrnRMSE:	0.8709,	TestRMSE:	0.9458
[20.64]	Epoch:	22,	TrnRMSE:	0.8657,	TestRMSE:	0.9424
[20.58]	Epoch:	23,	TrnRMSE:	0.8612,	TestRMSE:	0.9396
[20.56]	Epoch:	24,	TrnRMSE:	0.8571,	TestRMSE:	0.9375
[20.63]	Epoch:	25,	TrnRMSE:	0.8536,	TestRMSE:	0.9355
[21.15]	Epoch:	26,	TrnRMSE:	0.8504,	TestRMSE:	0.9334
[20.90]	Epoch:	27,	TrnRMSE:	0.8474,	TestRMSE:	0.9326
[20.94]	Epoch:	28,	TrnRMSE:	0.8447,	TestRMSE:	0.9308
[20.80]	Epoch:	29,	TrnRMSE:	0.8423,	TestRMSE:	0.9300
[20.51]	Epoch:	30,	TrnRMSE:	0.8399,	TestRMSE:	0.9292



Results (2)

■ Deep Matrix Factorization

[22.92]	Epoch:	1,	TrnRMSE:	1.6224,	TestRMSE:	1.0381
[23.06]	Epoch:	2,	TrnRMSE:	0.9955,	TestRMSE:	0.9624
[22.50]	Epoch:	3,	TrnRMSE:	0.9309,	TestRMSE:	0.9225
[22.31]	Epoch:	4,	TrnRMSE:	0.9105,	TestRMSE:	0.9154
[22.32]	Epoch:	5,	TrnRMSE:	0.9062,	TestRMSE:	0.9134
[22.05]	Epoch:	6,	TrnRMSE:	0.9042,	TestRMSE:	0.9116
[22.32]	Epoch:	7,	TrnRMSE:	0.9015,	TestRMSE:	0.9089
[22.33]	Epoch:	8,	TrnRMSE:	0.8967,	TestRMSE:	0.9053
[22.29]	Epoch:	9,	TrnRMSE:	0.8915,	TestRMSE:	0.9020
[22.27]	Epoch:	10,	TrnRMSE:	0.8874,	TestRMSE:	0.9004
[22.33]	Epoch:	11,	TrnRMSE:	0.8838,	TestRMSE:	0.8983
[22.11]	Epoch:	12,	TrnRMSE:	0.8805,	TestRMSE:	0.8968
[22.29]	Epoch:	13,	TrnRMSE:	0.8775,	TestRMSE:	0.8953
[22.33]	Epoch:	14,	TrnRMSE:	0.8744,	TestRMSE:	0.8949
[22.26]	Epoch:	15,	TrnRMSE:	0.8714,	TestRMSE:	0.8938
[22.30]	Epoch:	16,	TrnRMSE:	0.8686,	TestRMSE:	0.8937
[22.28]	Epoch:	17,	TrnRMSE:	0.8659,	TestRMSE:	0.8934
[22.08]	Epoch:	18,	TrnRMSE:	0.8634,	TestRMSE:	0.8939
[22.37]	Epoch:	19,	TrnRMSE:	0.8613,	TestRMSE:	0.8933
[22.37]	Epoch:	20,	TrnRMSE:	0.8592,	TestRMSE:	0.8931



What You Need to Know

- PyTorch
 - A Numpy-friendly deep learning framework
- Deep Matrix Factorization
 - A non-linear model
 - Wide range of parameters
 - A complex model



Questions?