

Jin-Soo Kim  
(jinsoo.kim@snu.ac.kr)

Systems Software &  
Architecture Lab.

Seoul National University

Jan. 3 – 14, 2022

*Python for Data Analytics*

# Data Preprocessing II



# Data Scaling

# Why Data Scaling?

- Features in a dataset can have very different scales
- Unscaled data can degrade the predictive performance of many machine learning algorithms
  - Many estimators assume that each feature takes values close to zero and all features vary on comparable scales
  - Metric-based and gradient-based estimators often assume approximately standardized data (normal distribution)
  - (cf.) Decision tree-based estimators are robust to arbitrary scaling of the data
- Unscaled data can slow down or even prevent the convergence of many gradient-based estimators

# Data Scaling

- Standard scaling:  $\rightarrow \tilde{x}_i \sim \text{Normal distribution } (\mu = 0, \sigma = 1)$   $\tilde{x}_i = \frac{x_i - \text{mean}(x)}{\text{std}(x)}$
- Min-Max Scaling:  $\rightarrow \tilde{x}_i \text{ in } [0, 1]$   $\tilde{x}_i = \frac{x_i - \min(x)}{\max(x) - \min(x)}$
- Max-Abs Scaling:  $\rightarrow |\tilde{x}_i| \leq 1$   $\tilde{x}_i = \frac{x_i}{\max(|x|)}$
- Robust Scaling:  $\rightarrow \text{Based on median and IQR}$   $\tilde{x}_i = \frac{x_i - \text{median}(x)}{Q3(x) - Q1(x)}$

## Data Scaling supported by SK-Learn

Scaling	Function	Class
Standard	<code>scale(x)</code>	<code>StandardScaler</code>
Min-Max	<code>minmax_scale()</code>	<code>MinMaxScaler</code>
Max-Abs	<code>maxabs_scale()</code>	<code>MaxAbsScaler</code>
Robust	<code>robust_scale()</code>	<code>RobustScaler</code>

# Standard Scaling

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing


df = pd.DataFrame({
    'x1': np.random.normal(0, 2, 10000),
    'x2': np.random.normal(5, 3, 10000),
    'x3': np.random.normal(-5, 5, 10000)
})

scaler = preprocessing.StandardScaler()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=['x1', 'x2', 'x3'])

sns.set_style('darkgrid')
_, (ax1, ax2) = plt.subplots(ncols=2, figsize=(6,5))
ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)
sns.kdeplot(df['x3'], ax=ax1)

ax2.set_title('After Scaling')
sns.kdeplot(scaled_df['x1'], ax=ax2)
sns.kdeplot(scaled_df['x2'], ax=ax2)
sns.kdeplot(scaled_df['x3'], ax=ax2)
```

$$\tilde{x}_i = \frac{x_i - \text{mean}(x)}{\text{std}(x)}$$



# Min-Max Scaling

- All values are mapped in the range [0, 1]
- Very sensitive to the presence of outliers

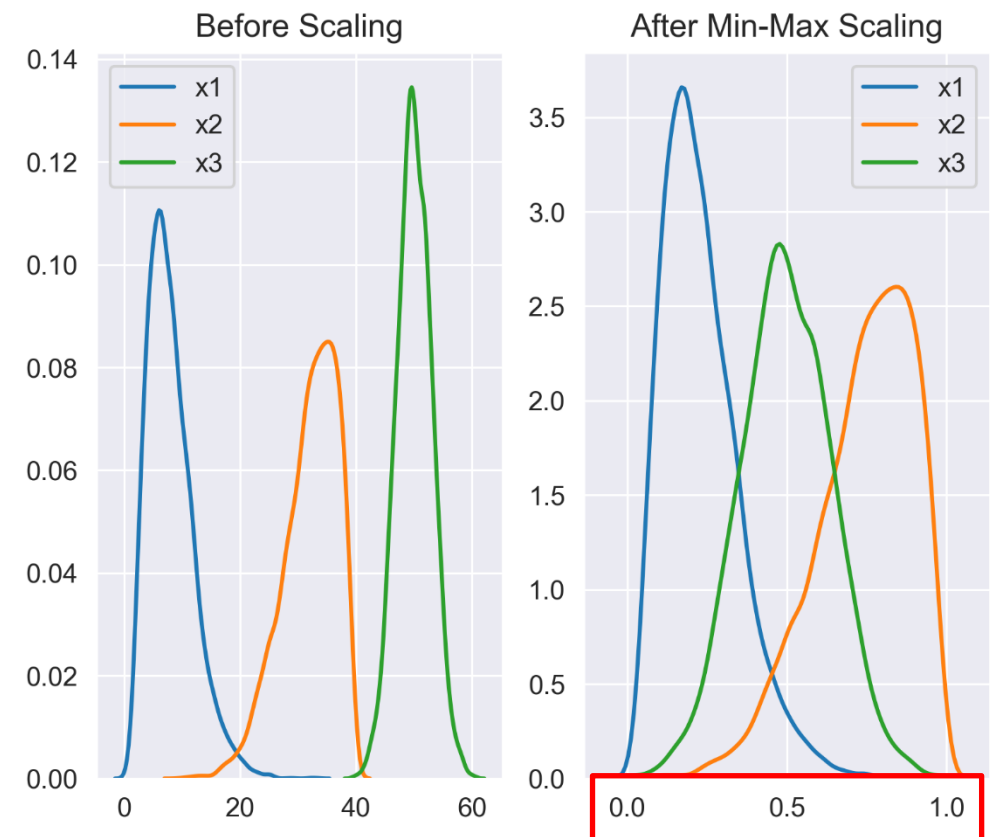
```
df = pd.DataFrame({
    'x1': np.random.chisquare(8, 10000),      # positive skew
    'x2': np.random.beta(8, 2, 10000)*40,    # negative skew
    'x3': np.random.normal(50, 3, 10000)     # no skew
})

scaler = preprocessing.MinMaxScaler()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=['x1', 'x2', 'x3'])

_, (ax1, ax2) = plt.subplots(ncols=2, figsize=(6,5))
ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)
sns.kdeplot(df['x3'], ax=ax1)

ax2.set_title('After Min-Max Scaling')
sns.kdeplot(scaled_df['x1'], ax=ax2)
sns.kdeplot(scaled_df['x2'], ax=ax2)
sns.kdeplot(scaled_df['x3'], ax=ax2)
```

$$\tilde{x}_i = \frac{x_i - \min(x)}{\max(x) - \min(x)}$$



# Max-Abs Scaling

- Doesn't change the shape of the distribution
- Also suffers from the presence of large outliers


$$\tilde{x}_i = \frac{x_i}{\max(|x|)}$$

```
df = pd.DataFrame({
    'x1': np.random.chisquare(8, 10000),      # positive skew
    'x2': np.random.beta(8, 2, 10000)*40,     # negative skew
    'x3': np.random.normal(50, 3, 10000)      # no skew
})

scaler = preprocessing.MaxAbsScaler()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=['x1', 'x2', 'x3'])

_, (ax1, ax2) = plt.subplots(ncols=2, figsize=(6,5))
ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)
sns.kdeplot(df['x3'], ax=ax1)

ax2.set_title('After Max-Abs Scaling')
sns.kdeplot(scaled_df['x1'], ax=ax2)
sns.kdeplot(scaled_df['x2'], ax=ax2)
sns.kdeplot(scaled_df['x3'], ax=ax2)
```



# Robust Scaling (I)

- Based on percentiles
- Not influenced by a few number of very large marginal outliers

$$\tilde{x}_i = \frac{x_i - \text{median}(x)}{Q3(x) - Q1(x)}$$

```
df = pd.DataFrame({
    # distribution with lower outliers
    'x1': np.hstack((np.random.normal(20,1,1000),
                     np.random.normal(1,1,25))),
    # distribution with upper outliers
    'x2': np.hstack((np.random.normal(30,1,1000),
                     np.random.normal(50,1,25)))
})

robust_scaler = preprocessing.RobustScaler()
robust_df = robust_scaler.fit_transform(df)
robust_df = pd.DataFrame(robust_df, columns=['x1', 'x2'])

minmax_scaler = preprocessing.MinMaxScaler()
minmax_df = minmax_scaler.fit_transform(df)
minmax_df = pd.DataFrame(minmax_df, columns=['x1', 'x2'])
```

```
_, (ax1, ax2, ax3) = plt.subplots(ncols=3, figsize=(9,5))
ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)

ax2.set_title('After Robust Scaling')
sns.kdeplot(robust_df['x1'], ax=ax2)
sns.kdeplot(robust_df['x2'], ax=ax2)

ax3.set_title('After Min-Max Scaling')
sns.kdeplot(minmax_df['x1'], ax=ax3)
sns.kdeplot(minmax_df['x2'], ax=ax3)
```




# Robust Scaling (2)

## ■ Min-Max Scaling

- Significantly affected by outliers

## ■ Robust Scaling


- Inliers are in  $[-2, 2]$
- Outliers still exist at the end of each distribution



# Data Standardization

# Data Skewness

- A measure of asymmetry of a distribution
  - Symmetrical (skewness = 0, e.g., normal distribution): mean == median == mode
  - Positive skew (skewness > 0): tail at right, mode < median < mean
  - Negative skew (skewness < 0): tail at left, mean < median < mode



# Measuring Data Skewness

- `df.skew([axis], [skipna], ...)`
  - Return unbiased skew over requested axis
  - `axis`: axis for the function to be applied on
  - `skipna`: if `True`, exclude null values when computing the result (default `True`)
- Meaning of skewness value
  - $-0.5 \leq \text{skewness} \leq 0.5$ : fairly symmetrical
  - $-1 < \text{skewness} < -0.5$  or  $0.5 < \text{skewness} < 1$ : moderately skewed
  - $\text{skewness} < -1$  or  $\text{skewness} > 1$ : highly skewed

# Handling Data Skewness

- Linear model performs better when the dataset follows normal distribution
- Dealing with positive skewness
  - Square root transformation ( $x$  to  $x^{1/2}$ )
  - Cube root transformation ( $x$  to  $x^{1/3}$ )
  - Log transformation ( $x$  to  $\log_2 x$ ,  $\log_e x$ ,  $\ln x$ , ...)
- Dealing with negative skewness
  - Square transformation ( $x^2$ )
  - Cube transformation ( $x^3$ )
  - Reflect the values and apply the methods used to reduce the positive skewness

# The Boston Housing Dataset

- Dataset for housing values in areas of Boston in 70's
- 506 rows, 14 columns (13 attributes + housing value)
- Available in the SK-Learn datasets

CRIM: 범죄율  
ZN: 25,000ft<sup>2</sup> 초과 거주지역 비율  
INDUS: 비소매상업지역 면적 비율  
CHAS: 찰스강 경계에 위치한 경우 1  
NOX: 일산화질소 농도  
AGE: 1940년 이전 건축된 주택 비율  
RM: 주택당 방 수  
RAD: 방사형 고속도로까지의 거리  
LSTAT: 인구 중 하위 계층 비율  
DIS: 직업 센터의 거리  
B: 인구 중 흑인 비율  
TAX: 재산세율  
PTRATIO: 학생/교사 비율  
**MEDV**: 주택 가격의 median (단위: \$1,000)

```
from sklearn import datasets
import pandas as pd
boston = datasets.load_boston()
df = pd.DataFrame(boston.data, columns=boston.feature_names)
df.head()
```


	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

# Skewness in Boston Housing Dataset


```
df.skew()
```

CRIM	5.223149
ZN	2.225666
INDUS	0.295022
CHAS	3.405904
NOX	0.729308
RM	0.403612
AGE	-0.598963
DIS	1.011781
RAD	1.004815
TAX	0.669956
PTRATIO	-0.802325
B	-2.890374
LSTAT	0.906460
dtype:	float64


```
plt.hist(df.CRIM, bins=20)
```




```
plt.hist(df.B, bins=20)
```



```
plt.hist(df.LSTAT, bins=20)
```



```
plt.hist(df.PTRATIO, bins=20)
```



# Transforming Data (I)

- `sklearn.preprocessing.scale(X, ...)`
  - Standardize a dataset along any axis (standard scaler)
  - Center to the zero mean and component wise scale to unit variance
  - `X`: the data to center and scale

```
from sklearn import preprocessing

df['LSTAT_log'] = preprocessing.scale(np.log(df['LSTAT']+1))
df['LSTAT_sqrt'] = preprocessing.scale(np.sqrt(df['LSTAT']+1))
df[['LSTAT', 'LSTAT_log', 'LSTAT_sqrt']].skew()
```

```
LSTAT      0.906460
LSTAT_log  -0.187195
LSTAT_sqrt  0.359606
dtype: float64
```




# Transforming Data (2)

Original data

```
import matplotlib.pyplot as plt


plt.hist(df['LSTAT'], bins=20)
plt.show()
```



Square root transformation

```
import matplotlib.pyplot as plt


plt.hist(df['LSTAT_sqrt'], bins=20)
plt.show()
```



Log transformation

```
import matplotlib.pyplot as plt

plt.hist(df['LSTAT_log'], bins=20)
plt.show()
```



# Sampling for Imbalanced Data

# Imbalanced Data

- A problem with classification where the classes are not represented equally
  - The model will be mostly tuned for the majority class
- Example:
  - A dataset with Class A : Class B = 9 : 1
  - The percentage of correct answers in the test dataset will also be 9 : 1
  - Even if a model classifies everything to Class A, it will have a 90% of accuracy
- Solutions: Balance data using sampling
  - Oversampling: increase the amount of minority class
  - Undersampling: use only part of majority class

# imbalanced-learn module

- A python package offering a number of re-sampling techniques
- Commonly used for datasets showing strong between-class imbalance
- Part of scikit-learn-contrib projects
- <https://github.com/scikit-learn-contrib/imbalanced-learn>
- Installation
  - pip install -U imbalanced-learn
  - conda install -c conda-forge imbalanced-learn

```
>>> import imblearn.under_sampling  
>>> import imblearn.over_sampling
```

# Creating Imbalanced Data

```
def plot(X, y):  
    plt.scatter(X[y==0, 0], X[y==0, 1], marker='x', label='Class 0')  
    plt.scatter(X[y==1, 0], X[y==1, 1], marker='o', label='Class 1')  
    plt.xlabel('X [0]')  
    plt.ylabel('X [1]')  
    plt.legend()
```

```
n0 = 450
```

```
n1 = 50
```


```
a = np.random.randn(n0, 2)*0.8 + 2 #  $N(2, 0.8)$ 
```

```
b = np.random.randn(n1, 2)*0.5 + 1 #  $N(1, 0.5)$ 
```

```
X = np.vstack([a, b])
```

```
y = np.hstack([np.zeros(n0), np.ones(n1)])
```

```
plot(X, y)
```



# Undersampling: RandomUnderSampler()

- Under-sample the majority class by randomly picking samples


```
from imblearn.under_sampling import RandomUnderSampler
```

```
X_samp, y_samp = RandomUnderSampler(random_state=0).fit_sample(X, y)
```

```
print(X_samp.shape, y_samp.shape)
```

```
plot(X_samp, y_samp)
```

```
(100, 2) (100,)
```



# Undersampling: EditedNearestNeighbours()

- Keep a sample if all or majority of the NN's belong to the same class


```
from imblearn.under_sampling import EditedNearestNeighbours

X_samp, y_samp = EditedNearestNeighbours(kind_sel='all', n_neighbors=10,
                                         random_state=0).fit_sample(X, y)

print(X_samp.shape, y_samp.shape)
plot(X_samp, y_samp)
```

(388, 2) (388,)

- n\_neighbors*: size of the neighbourhood to consider to compute the nearest neighbors
- kind\_sel*: 'all' (all have to agree to keep), 'mode' (majority vote to keep)



# Oversampling: RandomOverSampler()

- Over-sample the minority class by picking samples at random

```
from imblearn.over_sampling import RandomOverSampler
```


```
X_samp, y_samp = RandomOverSampler(random_state=0).fit_sample(X, y)
```

```
print(X_samp.shape, y_samp.shape)
```

```
plot(X_samp, y_samp)
```

(900, 2) (900,)

- Graph looks same, but the count has increased to 450





# Oversampling: SMOTE()

- A sample is created at a randomly selected point between a minority sample and its neighbor which is randomly selected among  $k$  neighbors



```
from imblearn.over_sampling import SMOTE
```

```
X_samp, y_samp = SMOTE(k_neighbors=3).fit_sample(X, y)
```

```
print(X_samp.shape, y_samp.shape)
```

```
plot(X_samp, y_samp)
```

(900, 2) (900,)



# Oversampling: ADASYN()

- For a minority sample dominated by majority class samples, more synthetic minority class samples are generated



```
from imblearn.over_sampling import ADASYN
```

```
X_samp, y_samp = ADASYN(n_neighbors=3, random_state=0).fit_sample(X, y)
```

```
print(X_samp.shape, y_samp.shape)
```

```
plot(X_samp, y_samp)
```

```
(908, 2) (908,)
```



Thank You!